

ECE491 Advanced compilers final report: Implementing a lazy functional language

Jonathan Lam

2022/05/12

Contents

1	Project overview	3
1.1	Motivation and overview	3
1.2	Commentary on the tutorial	3
1.3	Implementation setup	4
2	Background	5
2.1	Definition of a programming language	5
2.2	Implementation(s) of a programming language	5
2.3	The untyped λ -calculus	6
2.4	Functional programming	7
2.5	Terminology	8
3	The Core language	11
3.1	Syntax	11
3.2	Dynamic semantics	11
3.3	Base definitions	11
3.4	Sample programs	11
3.5	Lexer	11
3.6	Parser	11
3.7	Pretty-print utility	11
4	Template instantiation (TI) evaluator	12
4.1	The TI abstract machine	12
4.1.1	Instantiation	12
4.1.2	Unwinding	12
4.1.3	Updates and indirections for call-by-need	12
4.1.4	Memory model	12
4.2	Sample evaluations	12

5	G-Machine (GM) implementation	13
5.1	List of opcodes	13
5.2	Compilation schemes	14
5.2.1	Non-strict vs. strict compilation context	15
5.3	Evaluator	15
5.3.1	Transpilation to x86-64 assembly code	15
5.4	Sample compilations and evaluations	15
6	Future work	16
6.1	Garbage collection	16
6.2	Three-address machine and parallel G-machine	16
6.3	Particulars of the STG and Haskell's implementation	16
6.4	Compilation to native binaries	16
7	Conclusions	17
8	References	17

	Purely functional	Non-purely functional
Untyped	Core	Scheme ¹
Typed	Hazel ²	C ³

Table 1: Summary of programming language implementation projects

1 Project overview

1.1 Motivation and overview

This project was the semester-long project for the ECE491 independent study on advanced compilers. It is the implementation of a simple lazy (normal-order) evaluation functional language similar to Haskell. The work follows the tutorial, “Implementing functional languages: a tutorial” by Simon Peyton Jones at Microsoft Research [3]. This was published two years after Haskell 1.0 was defined in 1990 [1], to which SPJ was a major contributor.

This project is the culmination of my studies in programming languages and functional programming over the past two years. These studies include writing a C compiler in C (Spring 2021), a Scheme (LISP) interpreter in Scheme (Summer 2021), and this project, a Core interpreter and compiler in Haskell (Spring 2022). Additionally, my Master’s thesis (2021-2022 school year) was about updating the evaluation model of Hazel, a programming language implemented in OCaml. We summarize these languages in Table 1.

1.2 Commentary on the tutorial

The tutorial used for this project assumes a basic understanding of functional programming and a non-strict language such as Miranda or Haskell. For this report, we do not assume this and give a brief introduction to these ideas.

The text is largely in tutorial format, in that it provides much of the code for the reader to follow along with, but it is also in large part similar to a textbook. While much of the base code is given, much of the implementation is left in the form of non-trivial exercises to be completed using the provided theory. As the book progresses, less code is provided directly; rather, the semantics are given using the state transition notation, and the reader is asked to implement this in code.

The structure of the tutorial goes as follows: Chapter 1 introduces the Core language. Chapter 2 provides an evaluator implementation called the Template Instantiation (TI) evaluator, which we will describe in Section 4. Chapter 3 provides a compiled version of the TI evaluator, called the G-Machine⁴ (GM), which we will describe in Section 5. We note that Haskell officially uses a Spineless Tagless G-Machine (STG) [2].

¹Mostly functional

²Gradually-typed

³Weakly-typed

⁴“G” for “graph,” most likely.

Chapter 4 describes the Three-Address Machine compiled implementation, and Chapter 5 describes the Parallel G-Machine. The G-Machine (a stack-based machine) may be translated into a TAM representation. Chapter 6 introduces the λ -abstraction syntax using the λ -lifting program transformation. Chapters 4-6 were not covered for this independent study.

Rather than developing each implementation vertically (e.g., finishing the lexer, then the parser, then the intermediate representation, then exporting opcodes), the tutorial uses a horizontally incremental development method. In this style of development, we first complete a base implementation of only the core language features, and then incrementally add support for new language features. The tutorial calls these horizontal versions “marks,” i.e., “TI Mark 3” or “GM Mark 7.”

I greatly enjoyed the tutorial and found the exercises delightfully challenging and enlightening. My only complaints with the tutorial are that sometimes the exercises seemed to require knowledge from future marks, and that the names of certain concepts seem to be arbitrarily named⁵.

1.3 Implementation setup

My implementation is called `flc`, short for “fun(ctional) lazy compiler.” `flc` is written in Haskell, similar to the tutorial⁶. Instructions for use may be found on the GitHub’s README.

The implementation may be found at `jlam55555/fun-lazy-compiler`. A transpiler for the GM Mark 1 to x86-64 assembly code may be found at `jlam55555/flc-transpiler`.

⁵For example, the compilation schemes are named with arbitrary letters, and the “G” in “G-Machine” and the “I” in “Iseq” are never explained.

⁶The tutorial says it was written in Miranda, but I am not sure if this is correct. All of the code samples run successfully in Haskell, and I believe that Miranda’s syntax is somewhat different.

2 Background

This section provides brief background on programming language theory from a functional perspective.

2.1 Definition of a programming language

Interfaces are necessary for efficient and effective communication. A *programming language* serves as an interface between humans and computers. To rigorously work with a programming language, we define its *syntax* and *semantics*.

The syntax of a programming language describes the way valid expressions and programs are formed. It is specified using a *grammar*. For example, the grammar for the untyped λ -calculus Λ is shown in Figure 1.

The *semantics* of a programming language describes its behavior. The *static semantics* denotes the behavior of processes that happen prior to evaluation, such as type checking. The *dynamic semantics* describes the behavior of *evaluation*. Evaluation is the process of reducing an expression down to a *value*, or irreducible expression.

For this project, we define two languages: the Core language, and the abstract stack-based language of the G-machine. The semantics of the Core language is, like many functional languages, highly based on the λ -calculus. The semantics of the stack-based G-machine language is similar to other stack languages such as Forth.

2.2 Implementation(s) of a programming language

A programming language is a specification. There may be multiple *implementations* of a language. The specification and implementations are orthogonal. A single language may have multiple implementations with different performance characteristics and language support. For example, the Core language has four implementations proposed in the tutorial (TI, GM, TAM, PGM); other major languages such as the JVM, Python, or Haskell also have multiple implementations.

We typically classify an implementation into one of two broad classes: *interpreters* (a.k.a. *evaluators*) or *compilers*. The difference is that interpreters take a program’s source code and runs it “directly.” On the other hand, a compiler performs a two-stage process: compiling to a low-level representation, to be evaluated at some later time. Compiled representations tend to perform better and require simpler *runtimes*. The distinction between interpreters and compilers may not always be clear; in some cases, they form a spectrum delineated by the complexity of the runtime necessary for program evaluation.

In this project, we will examine two separate implementations of the Core language. The first implementation directly encodes the idea of graph template instantiation and reduction. Since all of the graph operations occur within Haskell code (a Haskell runtime), this is an interpreter. We improve this with the G-machine by compiling graph construction to a series of opcodes that may

$$e ::= \lambda x. e \mid x \mid e e$$

Figure 1: Grammar of Λ

run on a simpler stack-based machine, which may be implemented on many computer architectures with a simpler runtime⁷.

2.3 The untyped λ -calculus

The untyped λ -calculus Λ is a very basic model of universal computation⁸ proposed by Alonzo Church in the 1930's. The grammar of Λ is incredibly simple: it is given by Figure 1. We only have three expression forms in Λ : variables (bound by functions), λ -abstractions (functions of one variable x and return the expression e); and function application $e_1 e_2$, in which e_1 is in function position and e_2 is in argument position. Note that function application is represented using spaces, and parentheses are only necessary to specify order of operations.

The dynamic semantics are also very simple, illustrated by the rules shown in Figure 2 using a big-step operational semantics. Without diving deep into the syntax, this means that λ -abstractions evaluate to themselves, and the application of e_1 to e_2 involves (recursively) evaluating e_1 to a λ -abstraction $\lambda x. e'_1$; substituting e_2 for every instance of x in e'_1 ; and then recursively evaluating that result. This is a very simple way to think about function application, and it can be shown that any data or computation can be abstractly represented using this minimalistic representation⁹.

The λ -calculus serves as the logical and notational basis for much of functional programming, including Haskell and Core. For example, using space as the function application operator originates from the λ -calculus. Also, lazy evaluation and pure functions mesh very nicely with the more mathematical framework that the λ -calculus provides, as opposed to the statuful model of imperative programming languages.

However, since Λ is so minimalistic, it is not practical nor efficient. Usually we extend Λ with static typing, resulting in the simply-typed λ -calculus (STLC). In our case, we add a base type (integers) and structured data, and extend the grammar with `let` and `case` expressions, but do not implement static typing (which runs us the possibility of run-time type errors).

⁷Chapter 4 of the tutorial, which is not implemented for this project, uses a three-address machine, which may require an even simpler runtime on a three-address architecture such as x86-64.

⁸I.e., Λ can simulate a Turing machine.

⁹See Church notation.

$$\frac{}{\lambda x.e \Downarrow \lambda x.e} \qquad \frac{e_1 \Downarrow \lambda x.e'_1 \quad [e_2/x]e'_1 \Downarrow e}{e_1 \ e_2 \Downarrow e}$$

Figure 2: Dynamic semantics of Λ

2.4 Functional programming

Functional programming is a programming paradigm highly involved with function application, function composition, and first-class functions. It is a subclass of the declarative programming paradigm, which is concerned with pure expression-based computation. Declarative programming is often considered the complement of imperative programming, which may be characterized as programming with mutable state, side effects, or statements. Purely functional programming is a subset of functional programming that deals solely with mathematically-pure functions; non-pure languages may allow varying degrees of mutable state but typically encourage the use of pure functions. Miranda, Haskell, and Core are examples of non-strict (lazy) pure functional languages.

There are a number of memes¹⁰ among functional languages that may not be familiar to the imperative programmer. Some of these memes include (but are not limited to):

Spaces for function application This is a notational meme from Λ .

Immutability (i.e., lack of state) is an integral part of functional programming that allows for many optimizations such as structural sharing¹¹ or fast structural equality checking.

No side-effects Pure functional languages additionally disallow side-effects (stateful actions) such as printing or array-based operations, unless specially encapsulated. This may make debugging via printing difficult, so we need a special `Print` opcode in Core.

Curry-by-default functions and partial application In some functional languages, functions of multiple variables may be *partially applied*. For example, consider the following Haskell program.

```
sum x y = x + y
add1 x = sum 1
main = add1 3
```

This may seem to be an error, since `sum` is only applied to one argument in the definition of `add1`. However, this is allowable because `sum` will not be invoked until it is applied to the right number of arguments, which will happen after applying `add1` to 3. If we treat function application

¹⁰In the original sense of the word, not in the Internet sense of the word.

¹¹We implement structural sharing for integer nodes.

as substitution, this may be more clear: `add1 3` may be substituted to `(sum 1) 3` \equiv `sum 1 3`¹². Alternatively, we may think of any multi-arity function as being a series of nested λ -abstractions of one argument. The following is an equivalent definition of `sum`.

```
sum = \x . \y . x + y
```

Algebraic datatypes (ADTs) In Haskell, you can create type definitions such as the following:

```
data ListNode a = Nil | Cons a ListNode
data TreeNode = Leaf | Node Int TreeNode TreeNode
```

In this example, we see that we can very naturally describe linked lists using a recursive data structure with one of two forms: the base case `Nil`, or a `Cons` cell comprising a value and another (recursive) `ListNode`. This `ListNode` datatype is also parametrically polymorphic with type parameter `a`. We see a similar case for the (non-parametric, recursive) definition of `TreeNode`. ADTs allow us to naturally and safely express composite data and variant types, and language constructs such as pattern-matching and `case` expressions allow for safe deconstruction of structured data. In Core, we will implement a simpler-to-implement system of data constructors than user-defined data constructors as shown in the example above, but it is equally as powerful.

2.5 Terminology

The tutorial assumes a background in (lazy) functional programming. Instead, I assume a background in imperative programming and a basic background in programming language compilers. Below are some terms that may be useful for understanding Haskell, Core, and this project in general. These terms are fairly standard and can be easily looked up on Wikipedia or Stack Overflow, but the following are my own definitions. I thought I'd include them here to avoid confusion, since some of them have somewhat unassuming names.

Lazy evaluation A general term that describes evaluation that only occurs when the expression is used. For the purposes of this paper, it refers to *normal-order* evaluation, and more particularly *call-by-need* evaluation. Note that laziness (and eagerness) can refer to

Eager evaluation The opposite of lazy evaluation: evaluation that occurs when an expression is encountered. The eager analogue to normal-order evaluation is *applicative-order* evaluation.

Normal-order evaluation Lazy function application. In other words, the arguments to a function are not evaluated before the function body is

¹²These are equivalent due to the operator precedence and associativity of the infix function application operator (space).

evaluated. This is usually a feature only of pure functional languages such as Haskell or Core.

Applicative-order evaluation Eager function application. In other words, the arguments to a function are evaluated before the function body is evaluated. This is the way that most languages are implemented, such as C or OCaml.

Strict/non-strict More or less synonyms for eager/lazy. These terms are usually used when referring to Haskell’s normal-order evaluation. Haskell allows you to control strictness using special operators¹³. We explore controlling strictness with strict/non-strict compilation contexts in Section 5.2.1.

Call-by-value More or less synonymous with applicative-order evaluation.

Call-by-name A naïve form of normal-order evaluation. The unevaluated expression assigned to the function argument is lazily substituted into each instance of the argument, and evaluated when encountered.

Call-by-need A better form of normal-order evaluation, in which evaluation of arguments are memoized. Haskell and Core are call-by-need. This requires node updates and indirections as described in Section 4.1.3.

Supercombinator The top-level expression construct in Haskell and Core. Roughly equivalent to a function definition. All supercombinators are always in scope, including from the body of the supercombinator or before it is declared, allowing for easy simple and mutual recursion.

λ -abstraction Also known as λ -function, anonymous function, arrow function, function object, or simply function, in a language with first-class functions. In the λ -calculus, these typically take exactly one argument; however, in the case of Core or Haskell supercombinators or anonymous functions, these may take multiple arguments.

Function application Another name for function invocation, usually used in the context of functional programming. Note that functions may be *partially applied* in languages that support *currying-by-default*.

Currying A notational convenience in which a function of n arguments may be *partially-applied* by only calling it with $m < n$ arguments. This partially-applied expression itself is a (curried) function of $n - m$ arguments. When the function is applied to all n arguments, then the function is applied normally. This notation fits nicely into the interpretation of nested one-argument λ -abstractions, and is nicely handled by the TI and GM machines.

¹³https://wiki.haskell.org/Performance/Strictness#Explicit_strictness

Currying-by-default A property of a programming language in which functions of multiple arguments may be curried (partially applied). Haskell, OCaml, and Core are examples of such languages. In languages without currying-by-default, currying may be achieved using nested one-argument λ -abstractions.

Product types Composite datatypes, e.g., structs (a.k.a., objects, **records**, or **named tuples**) or tuples.

Sum types Disjunctive datatypes, notably *tagged unions*.

Algebraic datatypes (ADTs) Composite datatypes allow for arbitrary compositions of sum (tagged union) and product (tuple) types. These are a common construct available in many strongly-typed functional languages, such as Haskell, OCaml, and Rust. These are safely deconstructed using **case** expressions and pattern-matching. Data that is formed from ADTs are known as *structured data*.

Weak head normal form (WHNF) An expression is in WHNF if the outermost expression is a data constructor or a λ -abstraction. In other words, it defines what a fully-evaluated expression is in Core. Alternatively, we can think of it as the canonical representation of a datum in Core. (There are also normal and head normal forms, which are not relevant to Core.)

Scrutinee The expression to be matched in a **case** expression. In the case of the Core language, this must be structured data, but it usually is any pattern to match against.

Definiens The expression that a variable is bound to.

Opcode A.k.a. instruction. Simple statement in an imperative assembly-like language.

Runtime For the purposes of this project, usually refers to the run-time evaluator for a programming language implementation. As opposed to “runtime,” which is an adjective meaning “occurring at evaluation time.”

3 The Core language

3.1 Syntax

3.2 Dynamic semantics

TODO: function application and basic forms have the same dynamic semantics as the lambda calculus, except using an environment to store variable bindings rather than using substitution

3.3 Base definitions

TODO: standard representations (booleans, list constructors), standard definitions/prelude (e.g., if, simple combinators, standard representations)

3.4 Sample programs

TODO: show basic forms

 TODO: twice twice twice

 TODO: show algebraic datatypes

3.5 Lexer

TODO: very simple tokenization

3.6 Parser

TODO: build up LL(1) parser from useful subparsers

3.7 Pretty-print utility

TODO: constant-time append with indentation, only linear-time rendering

4 Template instantiation (TI) evaluator

TODO: only cover this briefly, relation to the GM machine

4.1 The TI abstract machine

TODO: process: instantiate and unwind

4.1.1 Instantiation

TODO: conditionals

4.1.2 Unwinding

4.1.3 Updates and indirections for call-by-need

4.1.4 Memory model

4.2 Sample evaluations

TODO: sample updating

TODO: sample lazy evaluation

TODO: tco

TODO: arithmetic

5 G-Machine (GM) implementation

The G-Machine is an abstract machine with a very similar idea to the TI machine with graph reduction and unwinding. The difference is that the graph creation, which is a complicated process performed by the instantiation step of the TI machine, is compiled into a series of opcodes on a stack-based machine. The unwinding step remains largely the same. The state of the machine is largely the same, with a stack, dump, and heap. Notably, this also includes an instruction queue (which replaces instantiation), and supercombinators are compiled to a sequence of instructions.

This stack-based machine is desirable for several efficiency reasons: it requires a simpler (smaller) runtime; it does not require the re-traversal of a supercombinator's body expression on each expression; and the runtime does not require an expression-level representation of the language.

5.1 List of opcodes

The list of opcodes of the abstract stack machine that the G-Machine compiles to is given below. For the sake of this document, the dynamic semantics of each opcode is only provided informally¹⁴. A precise definition of most of the opcodes are given in the tutorial¹⁵.

Pushglobal f Push the address of the global (supercombinator) f onto the stack¹⁶.

Pushint n Push the address of an integer node representing the integer n onto the stack. If such an integer node representing the integer n already exists, this may use that existing address; otherwise, this will allocate a new integer node n .

Push n Push the n -th address of the stack onto the stack.

Update n Set the node pointed to by the n -th address of the stack to an indirection node pointing to the top element of the stack, then pop one element from the stack.

Pop n Pop n elements from the stack.

Alloc n Push n invalid addresses onto the stack, or decrease the stack pointer by n ¹⁷.

¹⁴The syntax of this assembly-like language is trivial and need not be stated.

¹⁵Most of the time, the state transition(s) of an opcode are provided, and the implementation of the opcode in Haskell is left as an exercise.

¹⁶This is overloaded in Mark 6 for a technicality regarding unsaturated data constructors. This is Exercise 3.38.

¹⁷This is used for **letrec** expressions. These addresses initially point to an invalid node, and will be updated by the evaluation of the definiens.

Slide n “Slide” the top element of the stack up n slots. In other words, the $(n + 1)$ -th element of the stack will be replaced with the top of the stack, and then n elements will be popped.

Unwind Same as in the TI implementation Section 4.1.2.

Mkap Pop two elements from the stack, allocate a new application node (with the first popped element in function position, and the second popped element in argument position) and push the application node onto the stack.

Eval Strictly evaluate the element on the top of the stack to WHNF on a new stack (frame)¹⁸.

Add, Sub, Mul, Div, Neg Arithmetic binary and unary operators. Assumes the two elements on the top of the stack are evaluated to integers (otherwise will crash the program). Pops the top two elements and leaves a number node on the top of the stack representing the arithmetic result.

Eq, Ne, Lt, Le, Gt, Ge Relational binary operators. Assumes the two elements on the top of the stack are evaluated to booleans in the standard representation (otherwise will crash the program). Pops the top two elements and leaves a node on the top of the stack representing the boolean result in the standard representation.

Pack $t\ n$ Pops n addresses from the stack. Allocate and push a structured data node representing structured data with tag t and the n popped arguments.

Casejump rules The set *rules* is of the form $\{t_j \rightarrow [i_1, i_2, \dots]\}$: a mapping of tags t_j to code sequences $[i_1, i_2, \dots]$. Assumes the top of the stack is the scrutinee of a **case** expression, and thus a data node with tag t . Appends the instruction list of the appropriate rule onto the current instruction queue. Throws an error if the scrutinee is not a structured data node or if there is no match.

Split n Assumes the top of the stack is a structured data node. Pops it from the stack, and then pushes its n arguments onto the stack.

Print Output the top element on the stack. Assumes that it is a data node (integer or structured data node). Recursively generates additional **Print** and **Eval** opcodes for structured data.

5.2 Compilation schemes

TODO: list each compilation scheme, describe it, show its behavior on expression types

¹⁸This pushes a new stack on the dump and evaluates the node on the new stack. This is more or less equivalent to a function call in the conventional procedural ABIs.

5.2.1 Non-strict vs. strict compilation context

TODO: fit this in nicely with the previous section???

5.3 Evaluator

The evaluator implements the abstract stack machine. Each of the opcodes described in Section 5.1 is implemented as a separate function in the G-Machines's `Evaluator` module.

5.3.1 Transpilation to x86-64 assembly code

To illustrate the simplicity of the G-Machine's runtime, the stack machine of the Mark 1 of the G-Machine was implemented in NASM x86-64 assembly code. Each stack machine opcode was implemented using an assembly macro. Super-combinator node definitions are also implemented with a macro. This can be found on GitHub at [jlam55555/flc-transpiler](https://github.com/jlam55555/flc-transpiler).

I did not have time to implement later Marks of the G-Machine in assembly code, and I believe that it is a reasonable project for future work. However, the increased complexity of later Marks will likely be inconvenienced by small helper functions written in C.

5.4 Sample compilations and evaluations

TODO: this section

6 Future work

6.1 Garbage collection

For a functional language, garbage collection in the runtime is very important, since nodes are automatically allocated without control of the user. Garbage collection is covered in the last part of the TI implementation after Mark 6. It was not implemented for sake of time.

6.2 Three-address machine and parallel G-machine

These are two additional implementations of the compiler that are covered in future chapters of the tutorial. Both are very intriguing: the three-address machine is closer to the architecture of many modern CPU architectures, and inherently parallelized reduction is also an interesting idea.

6.3 Particulars of the STG and Haskell’s implementation

It will be interesting to study the intricacies of Haskell’s implementation. Haskell, despite being a high-level, lazy functional language, is known to be very fast to the point of bafflement¹⁹, with optimized GHC-compiled programs usually only 2 to 5 times slower than gcc-compiled programs²⁰.

We’ve seen that the graph-reduction mechanism in the TI machine is relatively intuitive, but there is significant complexity in the construction (instantiation) of graphs. The G-machine achieves some efficiency by compiling the construction of a graph to code. It will be interesting to see additional optimizations that allow GHC to achieve its famed performance. While this tutorial is meant to be an introductory tutorial on such machines, it should be formative to read documents specifically about the design of Haskell, such as [2].

6.4 Compilation to native binaries

We achieved compilation of the Mark 1 G-machine to native binaries using assembly macros and the NASM assembler, as described in Section 5.3.1. We leave for future work the same work for the Mark 6 G-machine, which is significantly more complicated.

However, it may be more fruitful to attempt this after reading the chapter on the three-address machine, since this chapter should describe the translation from the stack-based machine to a three-address machine. Transpiling the three-address machine to x86-64 opcodes should also be a simpler process.

¹⁹<https://stackoverflow.com/q/35027952>

²⁰<https://benchmarksgame-team.pages.debian.net/benchmarksgame/fastest/ghc-gcc.html>

7 Conclusions

This tutorial teaches how to efficiently evaluate a pure lazy functional language. First, an intuitive method involving lazy graph reduction called the Template Instantiation (TI) evaluator is introduced. The runtime of the TI comprises alternating graph creation (instantiation) and reduction (unwinding) processes. The efficiency of the runtime is improved by compiling the graph creation process into a series of opcodes in an stack machine, called the G-machine.

For this project, I was able to complete both base implementations and achieve the desired results. There is still much of the tutorial that has not been implemented, such as garbage collection, the λ -lifting program transformation, and other implementations such as the three-address machine and the parallel G-machine.

8 References

- [1] Paul Hudak et al. “A history of Haskell: being lazy with class”. In: *Proceedings of the third ACM SIGPLAN conference on History of programming languages*. 2007, pp. 12–1.
- [2] Simon L Peyton Jones. “Implementing lazy functional languages on stock hardware: the Spineless Tagless G-machine”. In: *Journal of functional programming* 2.2 (1992), pp. 127–202.
- [3] Simon L Peyton Jones and David R Lester. “Implementing functional languages: a tutorial”. In: *Department of Computer Science, University of Glasgow* (2000).